

Exemplar based approaches on faces

Thesis submitted in partial fulfillment
of the requirements for the degree of

MS by Research
in
Computer Science & Engineering

by

Mallikarjun B R
201307681

`kiran.raj@research.iiit.ac.in`



Center for Visual Information Technology
International Institute of Information Technology
Hyderabad - 500 032, India
February 2016

Copyright © Mallikarjun B R, 2016

All Rights Reserved

International Institute of Information Technology
Hyderabad, India

CERTIFICATE

It is certified that the work contained in this thesis, titled “Exemplar based approaches for faces” by Mallikarjun B R, has been carried out under my supervision and is not submitted elsewhere for a degree.

Date

Adviser: Prof. C V Jawahar

To My Family

Acknowledgments

Acknowledgments

Abstract

Multiplying two sparse matrices, denoted spmm , is a fundamental operation in linear algebra with several applications. Hence, efficient and scalable implementation of spmm has been a topic of immense research. Recent efforts are aimed at implementations on GPUs, multicore architectures, FPGAs, and such emerging computational platforms. Owing to the highly irregular nature of spmm , it is observed that GPUs and CPUs can offer comparable performance (Lee et al. [39]).

In this paper, we study CPU+GPU heterogeneous algorithms for spmm where the matrices exhibit a scale-free nature. Focusing on such matrices, we propose an algorithm that multiplies two sparse matrices exhibiting scale-free nature on a CPU+GPU heterogeneous platform.

Our experiments on a wide variety of real-world matrices from standard datasets show an average of 25% improvement over the best possible algorithm on a CPU+GPU heterogeneous platform. We show that our approach is both architecture-aware, and workload-aware.

The architectural trend towards heterogeneity has pushed heterogeneous computing to the fore of parallel computing research. Heterogeneous algorithms, often carefully handcrafted, are designed for several important problems from parallel computing such as sorting, graph algorithms, matrix computations, and the like. A majority of these algorithms follow a work partitioning approach where the input is divided into appropriate sized parts so that individual devices can process the right parts of the input. Such a division is done by means of thresholds. However, identifying the right value of the threshold is usually non-trivial and may require extensive empirical search. Such an extensive empirical search may potentially offset any gains accrued out of heterogeneous algorithms.

In this paper, we propose a simple and effective technique to identify the required thresholds in heterogeneous algorithms. Our technique is based on sampling and therefore can adapt to the algorithm used and the input instance. Our technique is generic in its applicability as we will demonstrate in this paper.

We validate our technique on two problems: finding the connected components of a graph and multiplying two scale-free sparse matrices. For these two problems, we show that using our method, we can find the required threshold that is $\pm 5\%$ and $\pm 7.5\%$ away from the best possible threshold, respectively. Along the way, we design a novel heterogeneous algorithm for sparse matrix multiplication when the matrices are scale-free in nature. This algorithm outperforms the existing best known algorithms for sparse matrix multiplication by 22% on average on a wide variety of matrices drawn from standard datasets.

Contents

Chapter	Page
1 Introduction	1
1.1 Relevance of Heterogeneous Computing	1
1.2 Sparse Matrix Computations	1
1.2.1 Sparse Matrix Sparse Matrix Multiplication	1
1.2.2 Sparse Matrix Dense Matrix Multiplication	1
1.3 Importance of Load Balancing	1
1.4 Contributions	1
2 Fiducial detection	2
2.1 Matrix Multiplication Formulation	2
2.1.1 The Row-Column Formulation	2
2.1.2 The Row-Row Formulation	2
2.1.2.1 Example of Row-Row Formulation	3
2.1.3 The Column-Row Formulation	3
2.1.3.1 Example of Column-Row Formulation	3
2.1.4 The Column-Column Formulation	4
2.1.4.1 Example of Column-Column Formulation	4
2.2 Sparse Matrix Storage Formats	5
2.2.1 Compressed Sparse Row (CSR) Format	5
2.2.2 Coordinate (COO) Format	5
2.3 Scale-free matrices	6
2.3.1 Real-world graphs	6
2.3.2 Synthetic graphs	6
2.4 Platform	6
2.4.1 CPU: OpenMP - Application Program Interface	6
2.4.2 GPU: CUDA - Application Program Interface	6
2.4.3 Heterogeneous Computing	6
2.4.4 Experimental Platform	6
2.5 Datasets	6
2.6 Related Work	6
2.6.1 Sparse Matrix Sparse Matrix Multiplication (SPMM)	6
2.6.2 Sparse Matrix Dense Matrix Multiplication (CSRMM)	6

3	Frontalization	7
3.1	Matrix Multiplication Formulation	7
3.1.1	The Row-Column Formulation	7
3.1.2	The Row-Row Formulation	7
3.1.2.1	Example of Row-Row Formulation	8
3.1.3	The Column-Row Formulation	8
3.1.3.1	Example of Column-Row Formulation	8
3.1.4	The Column-Column Formulation	9
3.1.4.1	Example of Column-Column Formulation	9
3.2	Sparse Matrix Storage Formats	10
3.2.1	Compressed Sparse Row (CSR) Format	10
3.2.2	Coordinate (COO) Format	10
3.3	Scale-free matrices	11
3.3.1	Real-world graphs	11
3.3.2	Synthetic graphs	11
3.4	Platform	11
3.4.1	CPU: OpenMP - Application Program Interface	11
3.4.2	GPU: CUDA - Application Program Interface	11
3.4.3	Heterogeneous Computing	11
3.4.4	Experimental Platform	11
3.5	Datasets	11
3.6	Related Work	11
3.6.1	Sparse Matrix Sparse Matrix Multiplication (SPMM)	11
3.6.2	Sparse Matrix Dense Matrix Multiplication (CSRMM)	11
4	Conclusions and Future Work	12
	Bibliography	14

List of Figures

Figure

Page

List of Tables

Table

Page

Chapter 1

Introduction

1.1 Relevance of Heterogeneous Computing

1.2 Sparse Matrix Computations

1.2.1 Sparse Matrix Sparse Matrix Multiplication

1.2.2 Sparse Matrix Dense Matrix Multiplication

1.3 Importance of Load Balancing

1.4 Contributions

Chapter 2

Fiducial detection

2.1 Abstract

Facial fiducial detection is a challenging problem for several reasons like varying pose, appearance, expression, partial occlusion and others. In the past, several approaches like mixture of trees [?], regression based methods [?], exemplar based methods [?] have been proposed to tackle this challenge.

In this paper, we propose an exemplar based approach to select the best solution from among outputs of regression and mixture of trees based algorithms (which we call candidate algorithms). We show that by using a very simple SIFT and HOG based descriptor, it is possible to identify the most accurate fiducial outputs from a set of results produced by candidate algorithms on any given test image. Our approach manifests as two algorithms, one based on optimizing an objective function with quadratic terms and the other based on simple kNN. Both algorithms take as input fiducial locations produced by running state-of-the-art candidate algorithms on an input image, and output accurate fiducials using a set of automatically selected exemplar images with annotations. Our surprising result is that in this case, a simple algorithm like kNN is able to take advantage of the seemingly huge complementarity of these candidate algorithms, better than optimization based algorithms.

We do extensive experiments on several datasets, and show that our approach outperforms state-of-the-art consistently. In some cases, we report as much as a 10% improvement in accuracy. We also extensively analyze each component of our approach, to illustrate its efficacy.

An implementation and extended technical report of our approach is available www.sites.google.com/site/wacv2016face

2.2 Related Work

2.3 Face Fiducial Detection

2.3.1 Formulation

2.3.2 Algorithm Outline

2.3.3 Exemplar Selection

2.3.4 Output selection by kNN

2.3.5 Output selection by Optimization

2.3.6 Implementation Details

2.4 Results

2.4.1 Qualitative Results

2.4.2 Experimental Analysis

2.5 References

Chapter 3

Frontalization

3.1 Abstract

3.2 Introduction

3.3 Related Work

3.4 Face Frontalization

3.4.1 Nearest exemplar selection

3.4.2 Triangulation and Transformation

3.4.3 Face Recognition

3.5 Experiments and Results

3.5.1 Exemplar Database

3.5.2 Comparision with Hassner et al

3.5.3 Quantitative Results

3.6 Conclusion

Chapter 4

Conclusions and Future Work

Conclusion

Related Publications

- *Kiran Raj Ramamoorthy, Dip Sankar Banerjee, Kannan Srinathan and Kishore Kothapalli, A Novel Heterogeneous Algorithm for Multiplying Scale-Free Sparse Matrices, IEEE - IPDPS, ASHES 2015.*
- *[Submitted] Hardhik Mallipeddi, Kiran Raj Ramamoorthy and Kishore Kothapalli, Nearly Balanced Work Partitioning via Sampling for Heterogeneous Algorithms, IEEE - 2016 .*

Bibliography

- [1] J. Alstott, E. Bullmore, and D. Plenz. Powerlaw: a Python package for analysis of heavy-tailed distributions. PLoS ONE 9(1): e85777, 2014.
- [2] Krste Asanovic, Ras Bodik, Bryan Christopher Catanzaro, Joseph James Gebis, Parry Husbands, Kurt Keutzer, David A. Patterson, William Lester Plishker, John Shalf, Samuel Webb Williams, and Katherine A. Yelick. The Landscape of Parallel Computing Research: A View from Berkeley. Technical Report No. UCB/EECS-2006-183, EECS Department, University of California, Berkeley.
- [3] D. A. Bader, V. Agarwal, K. Madduri. On the Design and Analysis of Irregular Algorithms on the Cell Processor: A Case Study of List Ranking. IPDPS, 2007.
- [4] D. A. Bader, and K. Madduri. GTgraph: A suite of synthetic graph generators. Available at <https://sdm.lbl.gov/kamesh/software/GTgraph/>
- [5] D. S. Banerjee, K. Kothapalli. Hybrid algorithms for list ranking and graph connected components. HiPC, 2011.
- [6] D. S. Banerjee, P. Sakurikar, and K. Kothapalli. Fast, scalable parallel comparison sort on hybrid multi-core architectures. In Proceedings of IPDPS Workshops, 2013.
- [7] D. S. Banerjee, S. Sharma, and K. Kothapalli. Work efficient parallel algorithms for large graph exploration. HiPC 2013.
- [8] M. Baskaran, and R. Bordawekar. Optimizing sparse matrix-vector multiplication on GPUs using compile-time and run-time strategies. Technical report, IBM Technical Report, 2008.
- [9] N. Bell, and M. Garland. Efficient sparse matrix-vector multiplication on CUDA. NVIDIA Technical Report, 2008
- [10] N. Bell, and M. Garland. Implementing sparse matrix-vector multiplication on throughput-oriented processors. In Proceedings of SC, 2009.
- [11] G. E. Blelloch, J. C. Hardwick, S. Chatterjee, J. Sipelstein, and M. Zagha. Implementation of a portable nested data-parallel language. In Proceedings of PPOPP, 1993.
- [12] G. E. Blelloch, M. A. Heroux, and M. Zagha. Segmented operations for sparse matrix computation on vector multiprocessors. Technical report, Pittsburgh, PA, USA, 1993.
- [13] M. Boyer, K. Skadron, S. Che, and N. Jayasena. Load Balancing in a Changing World: Dealing with Heterogeneity and Performance Variability. In Proceedings of ACM Computing Frontiers, 2013.

- [14] Aydin Buluc and John R. Gilbert. Challenges and advances in parallel sparse matrix-matrix multiplication. In The 37th International Conference on Parallel Processing (ICPP), 2008.
- [15] J. W. Choi, A. Singh, and R. W. Vuduc. Model-driven autotuning of sparse matrix-vector multiply on GPUs. In Proceedings of PPOPP, 2010.
- [16] J. Chhugani, N. Satish, Changkyu Kim, J. Sewall, P. Dubey. Fast and Efficient Graph Traversal Algorithm for CPUs: Maximizing Single-Node Efficiency. IPDPS, 2012.
- [17] T. Cormen, C. Leiserson, R. Rivest, and C. Stein. Introduction to algorithms. MIT Press, 2001
- [18] J.K. Cullum, and R.A. Willoughby. Lanczos Algorithms for Large Symmetric Eigenvalue Computations. Birkhauser, 1985.
- [19] T. A. Davis. Direct Methods for Sparse Linear Systems. Society for Industrial and Applied Mathematics, 2006.
- [20] M. deLorimier, and A. DeHon. Floating-point Sparse Matrix-Vector Multiply for FPGAs. FPGA, 2005.
- [21] J. W. Demmel. Applied Numerical Linear Algebra. Society for Industrial and Applied Mathematics, 1997.
- [22] M. Garland. Sparse matrix computations on many-core GPUs. In Proceedings of DAC, 2008.
- [23] R. A. V. D. Geijn, and J. Watts. SUMMA: Scalable universal matrix multiplication algorithm. Concurrency: Practice and Experience, 1997.
- [24] A. Gharaibeh, L. B. Costa, E. Santos-Neto, and M. Ripeanu. On graphs, GPUs, and blind dating: A workload to processor matchmaking quest. In Proceedings of IEEE IPDPS, 2013.
- [25] J. R. Gilbert, C. Moler, and R. Schreiber. Sparse matrices in Matlab: Design and implementation. SIAM Journal of Matrix Analysis and Applications, 1992.
- [26] G.H. Golub, and C.F. Van Loan. Matrix Computations. 2nd edition. Johns Hopkins University Press, Baltimore, 1989.
- [27] C. Gregg, M. Boyer, K. Hazelwood, and K. Skadron. Dynamic heterogeneous scheduling decisions using historical runtime data. in Workshop on Applications for Multi- and Many-Core Processors, 2011.
- [28] J. Greiner. A comparison of parallel algorithms for connected components. In Proceedings of SPAA, 1994.
- [29] D. Grewe and Michael F. P. OBoyle. A Static Task Partitioning Approach for Heterogeneous Systems Using openCL. In International Conference on Compiler Construction, 2011.
- [30] F. Gustavson. Two fast algorithms for sparse matrices: multiplication and permuted transposition. ACM Transactions on Mathematical Software (TOMS), 1978.
- [31] D. S. Hirschberg, A. K. Chandra, D. V. Sarwate. Computing connected components in parallel computers. Communications of the ACM, 1979.
- [32] Sungpack Hong, T. Oguntebi, K. Olukotun. Efficient Parallel Graph Exploration on Multi-Core CPU and GPU. PACT, 2011
- [33] S. Hong, N. C. Rodia, and K. Olukotun. On Fast Parallel Detection of Strongly Connected Components (SCC) in Small-World Graphs. In Proceedings of SC, 2013.

- [34] Q. Hou, K. Zhou, and B. Guo. SPAP: A programming language for heterogeneous many-core systems. Technical report, Zhejiang University, Graphics and Parallel Systems Lab, 2010.
- [35] S. Indarapu, M. Maramreddy, and K. Kothapalli. Architecture- and Workload-aware algorithms for Sparse Matrix-Vector Multiplication. ACM Compute, 2014.
- [36] K. Koer, I. Grasso, B. Cosenza, and T. Fahringer. An Automatic Input-Sensitive Approach for Heterogeneous Task Partitioning. In ACM ICS, 2013.
- [37] K. Kothapalli, S. Indarapu, S. Sharma, D. S. Banerjee, and R. Nigam. Workload Aware Algorithms for Heterogeneous Platforms, Under Submission, Available at <http://cstar.iiit.ac.in/kkishore/workqueue.pdf>, 2013
- [38] Zhiling Lan, Valerie E. Taylor, and Greg Bryan. A novel dynamic load balancing scheme for parallel systems. JPDC, 2002.
- [39] Victor W. Lee, Changkyu Kim, Jatin Chhugani, Michael Deisher, Daehyun Kim, Anthony D. Nguyen, Nadathur Satish, Mikhail Smelyanskiy, Srinivas Chennupaty, Per Hammarlund, Ronak Singhal, and Pradeep Dubey. Debunking the 100X GPU vs. CPU myth: an evaluation of throughput computing on CPU and GPU. In Proceedings of the 37th annual international symposium on Computer architecture. ISCA 2010.
- [40] J. Leskovec, J. Kleinberg, and C. Faloutsos. Graphs over time: densification laws, shrinking diameters and possible explanations. In Proceedings of ACM SIGKDD, 2005.
- [41] J. Leskovec, K. Lang, A. Dasgupta, and M. Mahoney. Community structure in large networks: Natural cluster sizes and the absence of large well-defined clusters. Internet Mathematics, 2009.
- [42] H. Ltaief, S. Tomov, R. Nath, and J. Dongarra. Hybrid Multicore Cholesky Factorization with Multiple GPU Accelerators. IEEE Transactions on Parallel and Distributed Computing, 2010.
- [43] C. K. Luk, S. Hong, and H. Kim. Qilin: Exploiting Parallelism on Heterogeneous Multiprocessors with Adaptive Mapping. In Proceedings of International Symposium on Microarchitecture (MICRO), 2009.
- [44] K. K. Matam, K. Kothapalli. Accelerating Sparse Matrix Vector Multiplication in Iterative Methods Using GPU. ICPP, 2011.
- [45] K. K. Matam, S. B. Indarapu, and K. Kothapalli. Sparse matrix-matrix multiplication on modern architectures. HiPC, 2012.
- [46] A. Monakov and A. Arutyun. Implementing Blocked Sparse Matrix-Vector Multiplication on NVIDIA GPUs. In Proceedings of SAMOS, 2009.
- [47] A. Monakov, A. Avetisyan, and A. Lokhmotov. Automatically tuning sparse matrix-vector multiplication for GPU Architectures. In Proceedings of HiPEAC, 2010.
- [48] R. Motwani, and P. Raghavan. Randomized Algorithms. Cambridge University Press, 2000.
- [49] S. C. Park, J. P. Draayer, and S. Q. Zheng. Fast sparse matrix multiplication. Computer Physics Communications, 1992.
- [50] Pawan Harish, and P. Narayanan. Accelerating Large Graph Algorithms on the GPU Using CUDA. In Proceedings of HiPC, 2007.

- [51] G. Penn. Efficient transitive closure of sparse matrices over closed semi-rings. *Theoretical Computer Science*, 2006.
- [52] W.H. Press, S.A. Teukolsky, W.T. Vetterling, and B.P. Flannery. *Numerical Recipes, The Art of Scientific Computing*. 2nd edition. Cambridge University Press, 1992.
- [53] M. O. Rabin and V. V. Vazirani. Maximum matchings in general graphs through randomization. *Journal of Algorithms*, 1989.
- [54] N. Satish, M. Harris, and M. Garland. Designing efficient sorting algorithms for many-core GPUs. In *Proceedings of IPDPS*, 2009.
- [55] D. P. Scarpazza, O. Villa, and F. Petrini. Efficient Breadth-First Search on the Cell/BE Processor. *IEEE Transactions on Parallel Distributed Systems*, 2008.
- [56] Shubhabrata Sengupta, Mark Harris, Yao Zhang, and John D. Owens. 2007. Scan primitives for GPU computing. In *Proceedings of GH*, 2007.
- [57] Y. Shan, W. Tianji, Y. Wang, B. Wang, Z. Wang, N. Xu, and H. Yang. FPGA and GPU implementation of large scale SpMV. *SASP*, 2010.
- [58] J. Shen, A. L. Varbanescu, P. Zou, Y. Lu, and H. Sips. Improving Performance by Matching Imbalanced Workloads with Heterogeneous Platforms. In *Proceedings of ACM ICS*, 2014.
- [59] Y. Shiloach, and U. Vishkin. An $O(\log n)$ parallel connectivity algorithm. *Journal of Algorithms*, 1982.
- [60] J. Siegel, O. Villa, S. Krishnamoorthy, A. Tumeo, Xiaoming Li. Efficient sparse matrix-matrix multiplication on heterogeneous high performance systems. *Cluster Computing Workshops and Posters*, IEEE 2010.
- [61] J. Soman, K. Kothapalli, and P. J. Narayanan. Some GPU Algorithms for Graph Connected Components and Spanning Tree, *Parallel Processing Letters*, 2010
- [62] Fengguang Song, Asim YarKhan, and Jack Dongarra. Dynamic task scheduling for linear algebra algorithms on distributed-memory multi-core systems. In *Proceedings of the Conference on High Performance Computing Networking, Storage and Analysis*, 2009.
- [63] P. D. Sulatycke, K. Ghose. Caching-efficient multi-threaded fast multiplication of sparse matrices. In *Proceedings of IPDPS*, 1998.
- [64] S. Tomov, J. Dongarra, and M. Baboulin. Towards dense linear algebra for hybrid GPU accelerated many-core systems. *Parallel Computing*, 2009.
- [65] V. Volkov, and J.W. Demmel. Benchmarking GPUs to tune dense linear algebra. In *Proceedings of SC*, 2008.
- [66] R. Vuduc. Automatic performance tuning of sparse matrix kernels. PhD thesis, UC Berkeley, USA, December 2003.
- [67] G. Wang and X. Ren. Power-efficient work distribution method for CPU-GPU heterogeneous system. In *Proceedings of ISPA*, 2010
- [68] Zheng Wei, and Joseph JaJa. Optimization of linked list prefix computations on multi-threaded GPUs using CUDA. In *Proceedings of IPDPS* 2010.

- [69] S. Williams, L. Oliker, R. Vuduc, J. Shalf, K. Yelick, and J. Demmel. Optimization of sparse matrix-vector multiplication on emerging multi-core platforms. In Proceedings of ACM SC, 2007.
- [70] I. Yamazaki, Tingxing Dong, S. Tomov, J. Dongarra. Tri-diagonalization of a Symmetric Dense Matrix on a GPU Cluster. IPDPSW, 2013
- [71] X. Yang, S. Parthasarathy, and P. Sadayappan, Fast Sparse Matrix-Vector Multiplication on GPUs: Implications for Graph Mining. In Proceedings of VLDB, 2011.
- [72] R. Yuster, and U. Zwick. Fast sparse matrix multiplication. ACM Transactions on Algorithms, 2005
- [73] L. Zhuo, V. K. Prasanna. Scalable and modular algorithms for floating-point matrix multiplication on FPGAs. In Proceedings of IPDPS, 2004.
- [74] U. Zwick. All pairs shortest paths using bridging sets and rectangular matrix multiplication. Journal of the ACM, 2002.
- [75] The Top 500 Super-computing sites. Available at <http://www.top500.org>
- [76] Timothy A. Davis, The University of Florida Sparse Matrix Collection, In Proceedings of NA DIGEST 1994.
- [77] Stanford Network Analysis Platform Dataset. <http://www.cise.ufl.edu/research/sparse/matrices/SNAP/>
- [78] openMP Application Programming Interface. <https://computing.llnl.gov/tutorials/openMP/>
- [79] CUDA: Compute unified device architecture programming guide. Technical report, NVIDIA.
- [80] Intel Math Kernel Library. <http://software.intel.com/en-us/articles/intel-mkl/>
- [81] nVidia cusp-library. <http://cusplibrary.github.io/>
- [82] nVidia cusparse Library. <http://developer.nvidia.com/cusparse/>